

✓ Regression Problem

- Explore feature correlation
- Simple Linear Regression
- Evaluate Model with MSE, MAE, R-square

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
columns = ['Overall Qual', 'Overall Cond', 'Gr Liv Area',
           'Central Air', 'Total Bsmt SF', 'SalePrice']

df = pd.read_csv('http://jse.amstat.org/v19n3/decock/AmesHousing.txt',
                 sep='\t',
                 usecols=columns)
```

```
df.head()
```

	Overall Qual	Overall Cond	Total Bsmt SF	Central Air	Gr Liv Area	SalePrice
0	6	5	1080.0	Y	1656	215000
1	5	6	882.0	Y	896	105000
2	6	6	1329.0	Y	1329	172000
3	7	5	2110.0	Y	2110	244000
4	5	5	928.0	Y	1629	189900

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
df.describe()
```

	Overall Qual	Overall Cond	Total Bsmt SF	Gr Liv Area	SalePrice
count	2930.000000	2930.000000	2929.000000	2930.000000	2930.000000
mean	6.094881	5.563140	1051.614544	1499.690444	180796.060068
std	1.411026	1.111537	440.615067	505.508887	79886.692357
min	1.000000	1.000000	0.000000	334.000000	12789.000000
25%	5.000000	5.000000	793.000000	1126.000000	129500.000000
50%	6.000000	5.000000	990.000000	1442.000000	160000.000000
75%	7.000000	6.000000	1302.000000	1742.750000	213500.000000
max	10.000000	9.000000	6110.000000	5642.000000	755000.000000

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2930 entries, 0 to 2929
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Overall Qual    2930 non-null  int64
1   Overall Cond    2930 non-null  int64
2   Total Bsmt SF   2929 non-null  float64
3   Central Air     2930 non-null  object
4   Gr Liv Area     2930 non-null  int64
```

```
5    SalePrice      2930 non-null    int64
dtypes: float64(1), int64(4), object(1)
memory usage: 137.5+ KB
```

```
# Encoding categorical feature
df['Central Air'] = df['Central Air'].map({'N': 0, 'Y': 1})
```

```
# Remove NaN
df = df.dropna(axis=0)
df.isnull().sum()
```

```
↔

```

	0
Overall Qual	0
Overall Cond	0
Total Bsmt SF	0
Central Air	0
Gr Liv Area	0
SalePrice	0

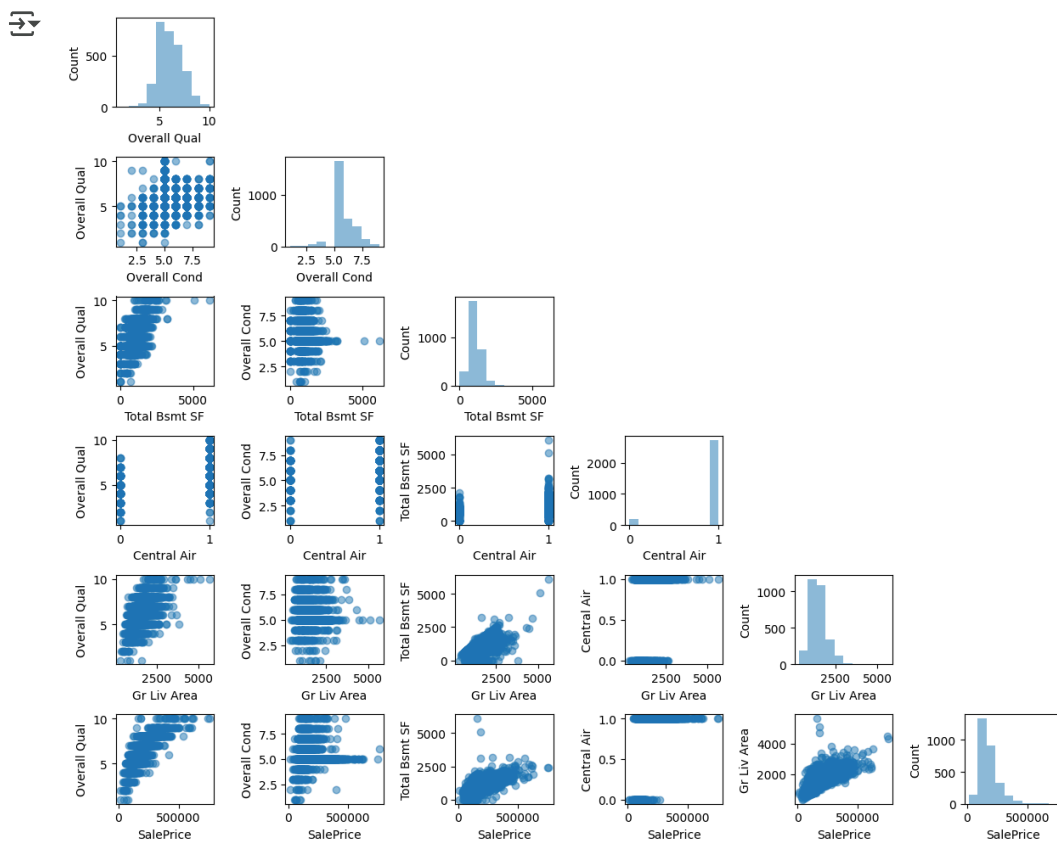
dtype: int64

✓ Visualizing Feature Correlation

```
import matplotlib.pyplot as plt
from mlxtend.plotting import scatterplotmatrix
```

```
# Use pairplot as an alternative
```

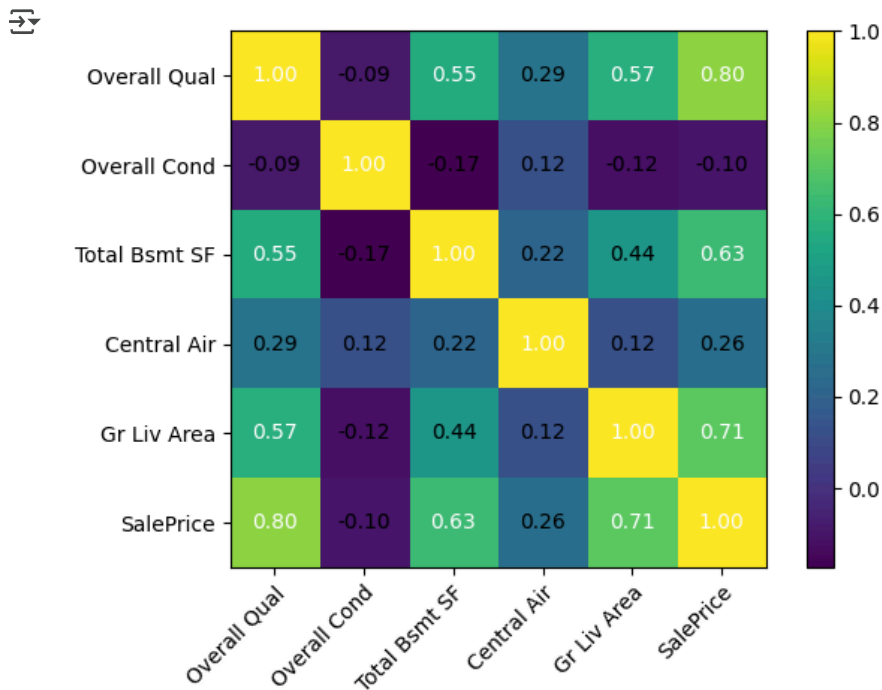
```
scatterplotmatrix(df.values, figsize=(12, 10),
                  names=df.columns, alpha=0.5)
plt.tight_layout()
```



```
import numpy as np
from mlxtend.plotting import heatmap

cm = np.corrcoef(df.values.T)
hm = heatmap(cm, row_names=df.columns, column_names=df.columns)

plt.tight_layout()
```



✓ Estimating the coefficient of a regression model

```
X = df[['Gr Liv Area']].values
y = df['SalePrice'].values
```

```
from sklearn.preprocessing import StandardScaler

sc_x = StandardScaler()
sc_y = StandardScaler()
X_std = sc_x.fit_transform(X)
y_std = sc_y.fit_transform(y[:, np.newaxis]).flatten()
```

```
from sklearn.linear_model import LinearRegression

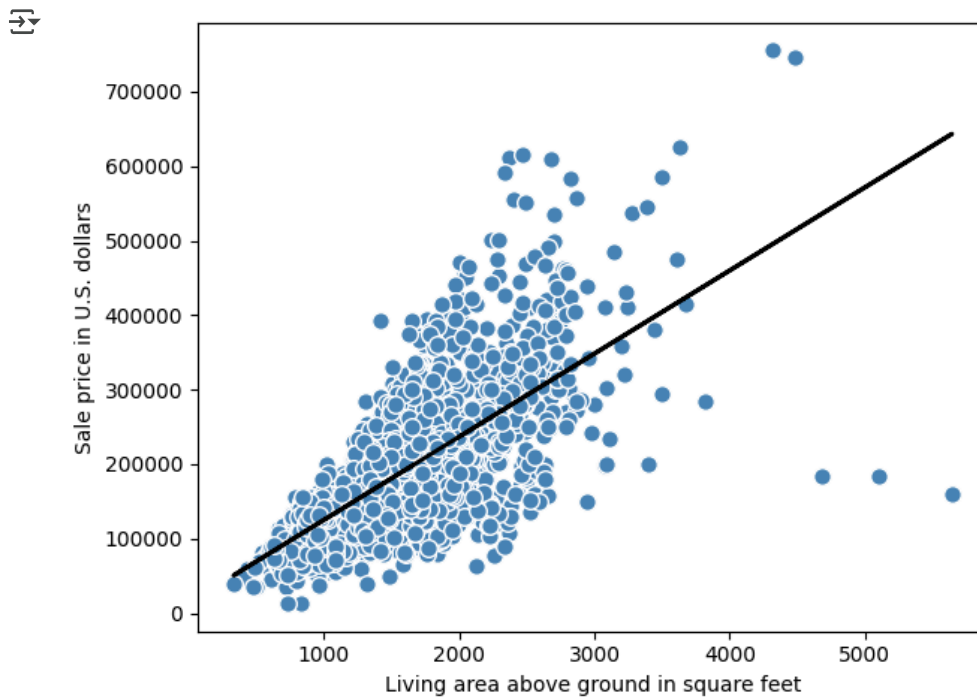
slr = LinearRegression()
slr.fit(X, y)
y_pred = slr.predict(X)
print(f'Slope: {slr.coef_[0]:.3f}')
print(f'Intercept: {slr.intercept_: .3f}')
```

```
➡ Slope: 111.666
Intercept: 13342.979
```

```
def lin_regplot(X, y, model):
    plt.scatter(X, y, c='steelblue', edgecolor='white', s=70)
    plt.plot(X, model.predict(X), color='black', lw=2)
    return

lin_regplot(X, y, slr)
plt.xlabel('Living area above ground in square feet')
plt.ylabel('Sale price in U.S. dollars')

plt.tight_layout()
```



✓ Fitting a robust regression model using RANSAC

```
from sklearn.linear_model import RANSACRegressor


ransac = RANSACRegressor(LinearRegression(),
                          max_trials=100, # default
                          min_samples=0.95,
                          loss='absolute_error', # default
                          residual_threshold=None, # default
                          random_state=123)

ransac.fit(X, y)

inlier_mask = ransac.inlier_mask_
outlier_mask = np.logical_not(inlier_mask)

line_X = np.arange(3, 10, 1)
line_y_ransac = ransac.predict(line_X[:, np.newaxis])
plt.scatter(X[inlier_mask], y[inlier_mask],
            c='steelblue', edgecolor='white',
            marker='o', label='Inliers')
plt.scatter(X[outlier_mask], y[outlier_mask],
            c='limegreen', edgecolor='white',
            marker='s', label='Outliers')
plt.plot(line_X, line_y_ransac, color='black', lw=2)
plt.xlabel('Living area above ground in square feet')
plt.ylabel('Sale price in U.S. dollars')
plt.legend(loc='upper left')

plt.tight_layout()
```



```
print(f'Slope: {ransac.estimator_.coef_[0]:.3f}')  
print(f'Intercept: {ransac.estimator_.intercept_: .3f}')
```

Slope: 106.348
Intercept: 20190.093

✓ Evaluating the performance of linear regression models

```
from sklearn.model_selection import train_test_split
```

```
target = 'SalePrice'  
features = df.columns[df.columns != target]
```

```
X = df[features].values  
y = df[target].values
```

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.3, random_state=123)
```

```
slr = LinearRegression()
```

```
slr.fit(X_train, y_train)  
y_train_pred = slr.predict(X_train)  
y_test_pred = slr.predict(X_test)
```

```
# Residual plot


x_max = np.max([np.max(y_train_pred), np.max(y_test_pred)])
x_min = np.min([np.min(y_train_pred), np.min(y_test_pred)])

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(7, 3), sharey=True)

ax1.scatter(y_test_pred, y_test_pred - y_test,
            c='limegreen', marker='s', edgecolor='white',
            label='Test data')
ax2.scatter(y_train_pred, y_train_pred - y_train,
            c='steelblue', marker='o', edgecolor='white',
            label='Training data')
ax1.set_ylabel('Residuals')

for ax in (ax1, ax2):
    ax.set_xlabel('Predicted values')
    ax.legend(loc='upper left')
    ax.hlines(y=0, xmin=x_min-100, xmax=x_max+100, color='black', lw=2)

plt.tight_layout()
```



```
from sklearn.metrics import mean_squared_error

mse_train = mean_squared_error(y_train, y_train_pred)
mse_test = mean_squared_error(y_test, y_test_pred)
print(f'MSE train: {mse_train:.2f}')
print(f'MSE test: {mse_test:.2f}')
```

```
➦ MSE train: 1497216245.85
  MSE test: 1516565821.00
```

```
from sklearn.metrics import mean_absolute_error

mae_train = mean_absolute_error(y_train, y_train_pred)
mae_test = mean_absolute_error(y_test, y_test_pred)
print(f'MAE train: {mae_train:.2f}')
print(f'MAE test: {mae_test:.2f}')
```

```
➦ MAE train: 25983.03
  MAE test: 24921.29
```

```
from sklearn.metrics import r2_score

r2_train = r2_score(y_train, y_train_pred)
r2_test = r2_score(y_test, y_test_pred)
print(f'R^2 train: {r2_train:.2f}')
print(f'R^2 test: {r2_test:.2f}')
```

```
➦ R^2 train: 0.77  
R^2 test: 0.75
```

✓ Apply regularization to regression

```
from sklearn.linear_model import Lasso  
  
lasso = Lasso(alpha=1.0)  
lasso.fit(X_train, y_train)  
y_train_pred = lasso.predict(X_train)  
y_test_pred = lasso.predict(X_test)  
print(lasso.coef_)
```

```
➦ [26251.38276394   804.70816337   41.94651964 11364.80761309  
   55.67855548]
```

```
train_mse = mean_squared_error(y_train, y_train_pred)  
test_mse = mean_squared_error(y_test, y_test_pred)  
print(f'MSE train: {train_mse:.3f}, test: {test_mse:.3f}')
```

```
train_r2 = r2_score(y_train, y_train_pred)  
test_r2 = r2_score(y_test, y_test_pred)  
print(f'R^2 train: {train_r2:.3f}, {test_r2:.3f}')
```

```
➦ MSE train: 1497216262.014, test: 1516576825.348  
R^2 train: 0.769, 0.752
```

```
from sklearn.linear_model import Lasso # L1
```

```
lasso = Lasso(alpha=1.0)
```

```
from sklearn.linear_model import Ridge # L2
```

```
ridge = Ridge(alpha=1.0)
```

```
from sklearn.linear_model import ElasticNet # L3
```

```
elanet = ElasticNet(alpha=1.0, l1_ratio=0.5)
```

✓ Dealing with nonlinear relationships using decision tree


```
from sklearn.tree import DecisionTreeRegressor
```

```
X = df[['Gr Liv Area']].values  
y = df['SalePrice'].values
```

```
tree = DecisionTreeRegressor(max_depth=3)  
tree.fit(X, y)  
sort_idx = X.flatten().argsort()
```

```
lin_regplot(X[sort_idx], y[sort_idx], tree)  
plt.xlabel('Living area above ground in square feet')  
plt.ylabel('Sale price in U.S. dollars')
```

```
plt.tight_layout()
```

✓ Your work

Use your own dataset or the **Boston Housing Dataset** below.

- Preprocess the dataset, e.g. drop null
- Explore the correlation of features with scatterplotmatrix or pairplot. Find if there is any highly correlated feature.
- Select the independent (low correlated) features as the input of your multiple regression model.
- Split dataset, standardize features and run simple linear regression model.
- Evaluate model with MSE, MAE and R-square.

Data Dictionary. <http://lib.stat.cmu.edu/datasets/boston>

Submit your notebook in PDF to BrightSpace by 4/20 11:59 pm.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
data_url = "http://lib.stat.cmu.edu/datasets/boston"
df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
```

```
even_rows = df.values[::2, :] # First part of each data point
odd_rows = df.values[1::2, :] # Second part with target
```


```
X = np.hstack([even_rows[:, :11], odd_rows[:, :2]]) # 13 features
y = odd_rows[:, 2] # MEDV target (3rd column of odd rows)
```

```
feature_names = ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE',
                 'DIS', 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT']
df = pd.DataFrame(X, columns=feature_names)
df['MEDV'] = y
```

```
print("Null values:", df.isnull().sum().sum())
#df = df.dropna()
#print("Null values after:", df.isnull().sum().sum())
```

Null values: 0

```
plt.figure(figsize=(12, 8))
sns.heatmap(df.corr(), annot=True, fmt=".2f", cmap='coolwarm')
plt.title('Feature Correlation Matrix')
plt.show()
```



```
corr_matrix = df.corr().abs()
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool))
to_drop = [column for column in upper.columns if any(upper[column] > 0.8)]
selected_features = [col for col in df.columns if col not in to_drop and col != 'MEDV']

print("Selected features:", selected_features)
```

Selected features: ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'PTRATIO', 'B', 'LSTAT']

```

X = df[selected_features].values
lr = LinearRegression()
lr.fit(X_train_scaled, y_train)

y_pred = lr.predict(X_test_scaled)

print("\nModel Evaluation:")
print(f"MSE: {mean_squared_error(y_test, y_pred):.2f}")
print(f"MAE: {mean_absolute_error(y_test, y_pred):.2f}")
print(f"R²: {r2_score(y_test, y_pred):.2f}")

```



Model Evaluation:
MSE: 22.47
MAE: 3.24
R²: 0.70

```

coefficients = pd.Series(lr.coef_, index=selected_features)
plt.figure(figsize=(10, 6))
coefficients.sort_values().plot.barh()
plt.title('Feature Importance (Linear Regression Coefficients)')
plt.show()

```

